

Name : Reena Qureshi

Reg No : 2023-BSE-052

Section: V-B

## LAB 12

### Task 1 — Organize Terraform code into separate files

Create the initial project structure:

```
mkdir -p ~/Lab12
```

```
@reenaquireshi ② /workspaces/cc_ReenQureshi_052 (main) $ mkdir -p ~/Lab12
@reenaquireshi ② /workspaces/cc_ReenQureshi_052 (main) $ cd ~/Lab12
@reenaquireshi ③ ~/Lab12 $
```

Create all required files:

```
touch main.tf variables.tf outputs.tf locals.tf terraform.tfvars entry-script.sh
```

Create **variables.tf** with the following content

```
@reenaquireshi ④ ~/Lab12 $ cat variables.tf
variable "vpc_cidr_block" {}
variable "subnet_cidr_block" {}
variable "availability_zone" {}
variable "env_prefix" {}
variable "instance_type" {}
variable "public_key" {}
variable "private_key" {}
@reenaquireshi ④ ~/Lab12 $
```

Create **outputs.tf** with the following content:

```
output "aws_instance_public_ip" {
    value = aws_instance.myapp-server.public_ip
}
```

```
tputs.tf@reenaquareshi ~ ~/Lab12 $ cat outputs.tf
output "aws_instance_public_ip" {
    value = aws_instance.myapp-server.public_ip
}
tputs.tf@reenaquireshi ~ ~/Lab12 $
```

Create `locals.tf` with the following content:

```
@reenaquireshi ~ ~/Lab12 $ cat locals.tf
locals {
    my_ip = "${chomp(data.http.my_ip.response_body)}/32"
}

data "http" "my_ip" {
    url = "https://icanhazip.com"
}
```

Create `terraform.tfvars` with the following content:

```
@reenaquireshi ~ ~/Lab12 $ cat terraform.tfvars
vpc_cidr_block = "10.0.0.0/16"
subnet_cidr_block = "10.0.10.0/24"
availability_zone = "us-east-1a"
env_prefix = "dev"
instance_type = "t3.micro"
public_key = "~/.ssh/id_ed25519.pub"
private_key = "~/.ssh/id_ed25519"
```

Create `main.tf` with the following content:

```
> EOF
@reenaquireshi ② ~/Lab12 $ cat main.tf
provider "aws" {
  region = "us-east-1"
}

resource "aws_vpc" "myapp_vpc" {
  cidr_block = var.vpc_cidr_block
  tags = {
    Name = "${var.env_prefix}-vpc"
  }
}

resource "aws_subnet" "myapp_subnet_1" {
  vpc_id      = aws_vpc.myapp_vpc.id
  cidr_block = var.subnet_cidr_block
  availability_zone = var.availability_zone
  tags = {
    Name = "${var.env_prefix}-subnet-1"
  }
}

resource "aws_default_route_table" "main_rt" {
  default_route_table_id = aws_vpc.myapp_vpc.default_route_table_id
  route {
    cidr_block = "0.0.0.0/0"
```

Create entry-script.sh with the following content:

```
└─
@reenaquireshi ② ~/Lab12 $ cat entry-script.sh
#!/bin/bash
set -e
yum update -y
yum install -y nginx
systemctl start nginx
systemctl enable nginx
@reenaquireshi ② ~/Lab12 $ chmod +x entry-script.sh
@reenaquireshi ② ~/Lab12 $ ssh-keygen -t ed25519 -f ~/.ssh/id_ed25519
~/.ssh/@reenaquireshi ② ~/Lab12 $ ls -la ~/.ssh/
total 24
```

Initialize Terraform:

```
@reenaquireshi ② ~/Lab12 $ terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Finding latest version of hashicorp/http...
- Installing hashicorp/aws v6.27.0...
- Installed hashicorp/aws v6.27.0 (signed by HashiCorp)
- Installing hashicorp/http v3.5.0...
- Installed hashicorp/http v3.5.0 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

Apply the configuration:

```
@reenaquireshi ② ~/Lab12 $ @reenaquireshi ② ~/Lab12 $ terraform apply -auto-approve
data.http.my_ip: Reading...
data.http.my_ip: Read complete after 1s [id=https://icanhazip.com]
aws_key_pair.ssh-key: Refreshing state... [id=serverkey]
aws_vpc.myapp_vpc: Refreshing state... [id=vpc-037a2b5f96140533f]
aws_subnet.myapp_subnet_1: Refreshing state... [id=subnet-01a103908f7624a14]
aws_internet_gateway.myapp_igw: Refreshing state... [id=igw-0820c6d8fc5e74bf0]
aws_default_security_group.default_sg: Refreshing state... [id=sg-0e28cdb94003274fe]
aws_default_route_table.main_rt: Refreshing state... [id=rtb-028ce09473ca6d6c3]

Terraform used the selected providers to generate the following execution plan. Resource actions
following symbols:
+ create
```

Display the output:

```
+ ipa...
aws_instance public_ip = "3.231.50.188"
```

Test nginx in browser:

Open browser and navigate to <http://<public-ip>>

---

## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org).  
Commercial support is available at [nginx.com](http://nginx.com).

*Thank you for using nginx.*

Destroy:

```
- vpc_id          = "vpc-0a9005120b91ce98c" -> null
  # (1 unchanged attribute hidden)
}

Plan: 0 to add, 0 to change, 3 to destroy.

Changes to Outputs:
- aws_instance_public_ip = "52.3.233.109" -> null
- ssh_command           = "ssh -i ~/.ssh/id_rsa_lab12 ubuntu@52.3.233.109" ->

Do you really want to destroy all resources?
Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

aws_instance.myapp-server: Destroying... [id=i-0372abda970304c73]
aws_instance.myapp-server: Still destroying... [id=i-0372abda970304c73, 00m10s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-0372abda970304c73, 00m20s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-0372abda970304c73, 00m30s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-0372abda970304c73, 00m40s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-0372abda970304c73, 00m50s elapsed]
aws_instance.myapp-server: Destruction complete after 52s
aws_key_pair.ssh-key: Destroying... [id=dev-serverkey-lab12]
aws_security_group.web_sg: Destroying... [id=sg-0db2d26bea296f02f]
aws_key_pair.ssh-key: Destruction complete after 1s
aws_security_group.web_sg: Destruction complete after 2s

destroy complete! Resources: 3 destroyed
```

## Task 2 — Use remote-exec provisioner

Modify the `aws_instance` resource in `main.tf` to use `remote-exec` provisioner:

Apply the configuration:

```

30
curl http://$IP@reenaquareeshi ~/Lab12 $ terraform apply -auto-approve

data.http.my_ip: Reading...
data.http.my_ip: Read complete after 0s [id=https://icanhazip.com]
data.aws_vpc.default: Reading...
aws_key_pair.ssh-key: Refreshing state... [id=serverkey]
data.aws_vpc.default: Read complete after 2s [id=vpc-0a9005120b91ce98c]
data.aws_subnet.default: Reading...
aws_security_group.default_sg: Refreshing state... [id=sg-010af2dd39e1db8db]
data.aws_subnet.default: Read complete after 0s [id=subnet-00447af4668f7212a]
aws_instance.myapp-server: Refreshing state... [id=i-06a42a3d0b7f34cf1]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with
following symbols:
  ~ update in-place

Terraform will perform the following actions:

# aws_instance.myapp-server will be updated in-place
~ resource "aws_instance" "myapp-server" {
    id                      = "i-06a42a3d0b7f34cf1"
    ~ public_dns              = "ec2-54-236-40-82.compute-1.amazonaws.com" -> (known after apply)
    ~ public_ip               = "54.236.40.82" -> (known after apply)
    tags                    = {
        "Name" = "dev-ec2-instance"
    }
    - user_data
      #!/bin/bash
      set -e
}

```

Display the output:

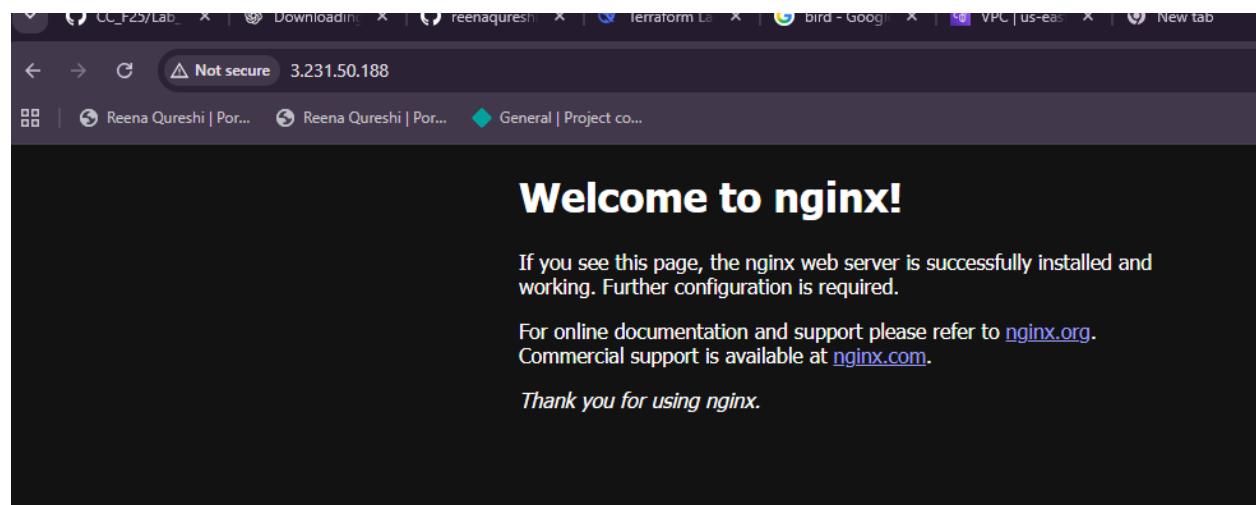
`terraform output`

```

aws_instance_public_ip = "3.231.50.188"

```

Test nginx in browser:



## Task 3 — Use file and local-exec provisioners

Modify the `aws_instance` resource in `main.tf` to include all three provisoners

```
sg"
  description = "Defa@reenaquireshi ~ ~/Lab12 $ cat > main.tf << 'EOF'
provider "aws" {
  region = "us-east-1"
}
#
# Use default VPC
data "aws_vpc" "default" {
  default = true
}
#
# Use default subnet
data "aws_subnet" "default" {
  vpc_id          = data.aws_vpc.default.id
  availability_zone = var.availability_zone
  default_for_az   = true
}
#
# Security Group
resource "aws_security_group" "default_sg" {
  name      = "${var.env_prefix}-default-sg"
  description = "Default security group for lab"
  vpc_id      = data.aws_vpc.default.id

  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
```

Apply the configuration:

```
aws_instance.myapp-server (remote-exec): Created symlink from /etc/systemd/system/multi-user.target.wants/nginx.service to /usr/lib/systemd/system/nginx.service.
aws_instance.myapp-server (remote-exec): Nginx installation complete from uploaded script!
aws_instance.myapp-server (remote-exec): • nginx.service - The nginx HTTP and reverse proxy server
aws_instance.myapp-server (remote-exec):     Loaded: loaded (/usr/lib/systemd/system/nginx.service; disabled)
aws_instance.myapp-server (remote-exec):     Active: active (running) since Sun 2026-01-04 10:35:28 UTC
aws_instance.myapp-server (remote-exec):       Main PID: 5699 (nginx)
aws_instance.myapp-server (remote-exec):         CGroup: /system.slice/nginx.service
aws_instance.myapp-server (remote-exec):             ├─5699 nginx: master process...
aws_instance.myapp-server (remote-exec):             ├─5701 nginx: worker process...
aws_instance.myapp-server (remote-exec):             ├─5702 nginx: worker process...
aws_instance.myapp-server (remote-exec): Jan 04 10:35:28 ip-172-31-2-214.ec2.internal systemd[1]: Starting The nginx HTTP and reverse proxy server...
aws_instance.myapp-server (remote-exec): Jan 04 10:35:28 ip-172-31-2-214.ec2.internal nginx[5699]: nginx/1.27.3 (Ubuntu)
aws_instance.myapp-server (remote-exec): Jan 04 10:35:28 ip-172-31-2-214.ec2.internal nginx[5699]: 2026/01/04 10:35:28 [info] 5699#5699: start of worker processes
aws_instance.myapp-server (remote-exec): Jan 04 10:35:28 ip-172-31-2-214.ec2.internal nginx[5699]: 2026/01/04 10:35:28 [info] 5699#5699: worker process 5701 started
aws_instance.myapp-server (remote-exec): Hint: Some lines were ellipsized, use -l to show more.
aws_instance.myapp-server: Provisioning with 'local-exec'...
aws_instance.myapp-server (local-exec): Executing: ["./bin/sh" "-c" "echo \"Instance id=i-064d1a62869946103 has been created at $(date)\" >> instance_creation.log\necho \"SSH command: ssh user@18.209.240.153\" >> instance_creation.log\n"]
aws_instance.myapp-server: Creation complete after 58s [id=i-064d1a62869946103]

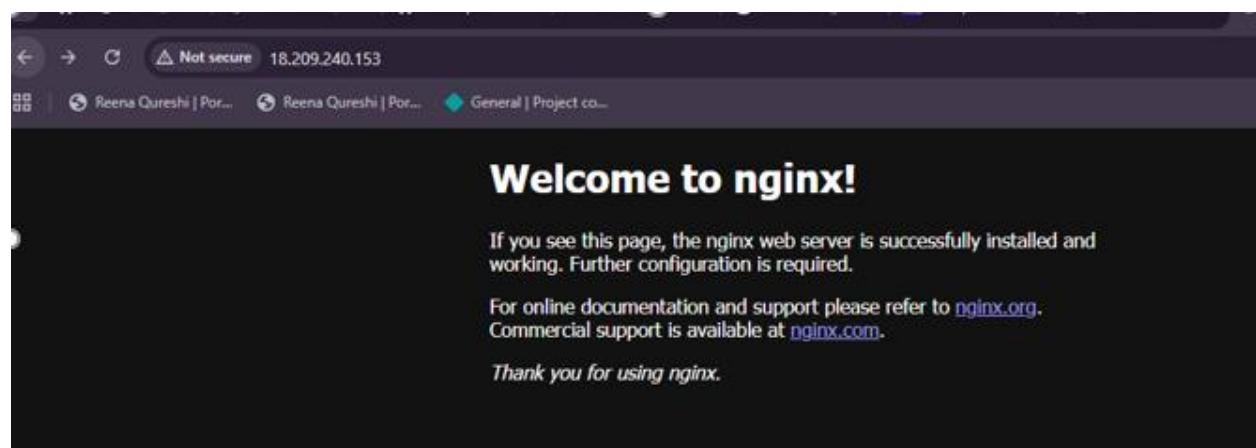
apply complete! Resources: 3 added, 0 changed, 0 destroyed.
```

Display the output:

terraform output

```
aws_instance_public_ip = "18.209.240.153"
```

Test nginx in browser:



Destroy the resources:

```
Plan: 0 to add, 0 to change, 3 to destroy.

Changes to Outputs:
- aws_instance_public_ip = "18.209.240.153" -> null
aws_instance.myapp-server: Destroying... [id=i-064d1a62869946103]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 00m10s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 00m20s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 00m30s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 00m40s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 00m50s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 01m00s elapsed]
aws_instance.myapp-server: Still destroying... [id=i-064d1a62869946103, 01m10s elapsed]
aws_instance.myapp-server: Destruction complete after 1m12s
aws_key_pair.ssh-key: Destroying... [id=serverkey]
aws_security_group.default_sg: Destroying... [id=sg-0551302a079993d3f]
aws_key_pair.ssh-key: Destruction complete after 0s
aws_security_group.default_sg: Destruction complete after 2s
```

## Task 4 — Create Terraform modules (subnet module)

In this task, you will create a reusable subnet module to organize your infrastructure code better.

Create the module directory structure:

```
mkdir -p modules/subnet
```

```
touch modules/subnet/main.tf
```

```
touch modules/subnet/variables.tf
```

```
touch modules/subnet/outputs.tf
```

```
greenaquireshi  ~/Lab12 $ tree
.
+-- entry-script.sh
+-- instance_creation.log
+-- locals.tf
+-- main.tf
+-- main.tf.task1
+-- modules
|   +-- subnet
|   |   |-- main.tf
|   |   |-- outputs.tf
|   |   `-- variables.tf
+-- outputs.tf
+-- screenshots
|   |-- task3_main_tf_all_provisioners.txt
|   |-- task3_main_tf_restored.txt
|   |-- task3_nginx_browser.txt
|   |-- task3_terraform_apply.txt
|   |-- task3_terraform_destroy.txt
|   `-- task3_terraform_output.txt
+-- terraform.tfstate
+-- terraform.tfstate.backup
+-- terraform.tfvars
+-- test_url.txt
+-- variables.tf

4 directories, 20 files
6. Create modules/subnet/main.tf :
```

Create modules/subnet/variables.tf:

```
variable "vpc_id" {}
variable "subnet_cidr_block" {}
variable "availability_zone" {}
variable "env_prefix" {}
variable "default_route_table_id" {}
```

Create modules/subnet/outputs.tf

```
greenaquireshi  ~/Lab12 $ cat modules/subnet/main.tf
data "aws_vpc" "default" {
  default = true
}

data "aws_subnet" "default" {
  vpc_id      = data.aws_vpc.default.id
  availability_zone = var.availability_zone
  default_for_az    = true
}
```

Initialize Terraform to download the module:

```
terraform init
```

```
@reenaquireshi ② ~/Lab12 $ terraform init
Initializing the backend...
Initializing modules...
- myapp-subnet in modules/subnet
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Reusing previous version of hashicorp/http from the dependency lock file
- Using previously-installed hashicorp/aws v6.27.0
- Using previously-installed hashicorp/http v3.5.0

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

Configure it:

```
+ vpc_id          = vpc-0a9005120b91ce98c
}

Plan: 3 to add, 0 to change, 0 to destroy.

Changes to Outputs:
+ aws_instance_public_ip = (known after apply)
aws_key_pair.ssh-key: Creating...
aws_security_group.default_sg: Creating...
aws_key_pair.ssh-key: Creation complete after 1s [id=serverkey]
aws_security_group.default_sg: Creation complete after 5s [id=sg-0d784d473193227ab]
aws_instance.myapp-server: Creating...
aws_instance.myapp-server: Still creating... [00m10s elapsed]
aws_instance.myapp-server: Creation complete after 15s [id=i-035202f3bcf507518]

Apply complete! Resources: 3 added, 0 changed, 0 destroyed.

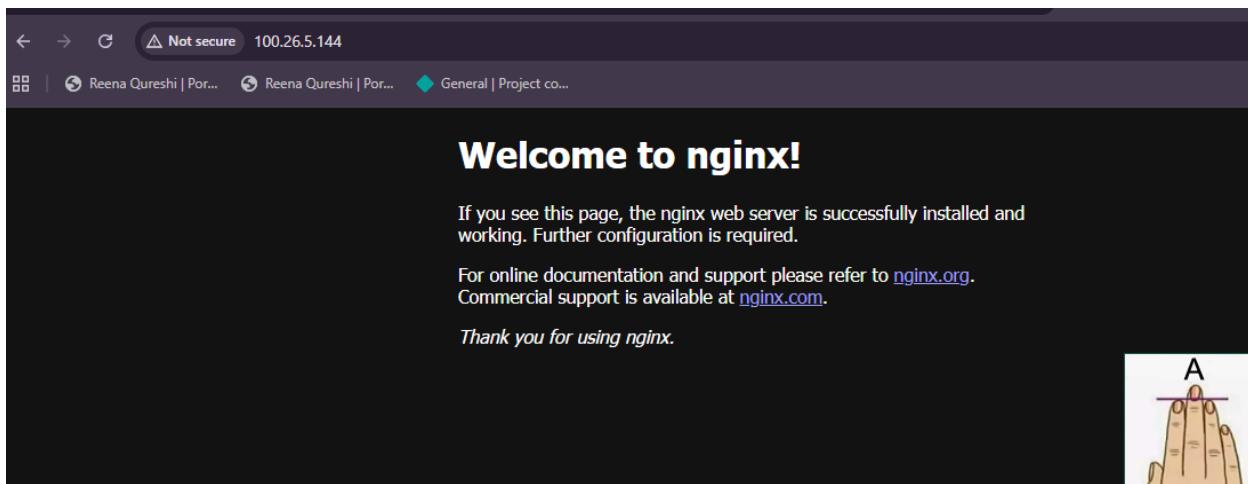
Outputs:

aws_instance public ip = "100.26.5.144"
```

OutPut:

```
outputs:
aws_instance_public_ip = "100.26.5.144"
@reenaquireshi ② ~/Lab12 $ client_loop: send disconnect: Connection abor
shell closed: exit status 255
```

BROWSER:



## Task 5 — Create webserver module

In this task, you will create a reusable webserver module for EC2 instances.

Create the webserver module directory structure:

```
mkdir -p modules/webserver  
touch modules/webserver/main.tf  
touch modules/webserver/variables.tf  
touch modules/webserver/outputs.tf
```

Modify the root `main.tf`:

Remove the security group, ke

```
@reenaquireshi ~ ~/Lab12 $ # Update root main.tf to use webserver module
@reenaquireshi ~ ~/Lab12 $ cat > main.tf << 'EOF'
> provider "aws" {
>   region = "us-east-1"
> }
>
> module "myapp-subnet" {
>   source = "./modules/subnet"
>   availability_zone = var.availability_zone
>   env_prefix = var.env_prefix
> }
>
> module "myapp-webserver" {
>   source = "./modules/webserver"
>   env_prefix = var.env_prefix
>   instance_type = var.instance_type
>   availability_zone = var.availability_zone
>   public_key = var.public_key
>   my_ip = local.my_ip
>   vpc_id = module.myapp-subnet.vpc_id
>   subnet_id = module.myapp-subnet.subnet.id
>   script_path = "./entry-script.sh"
>   instance_suffix = "0"
> }
> EOF
```

## Configuring:

```
git Select Command Prompt - git codespace ssh -c !unify-bassoon=Jg7q7jwqpg0n9px
+ owner_id          = (known after apply)
+ region            = "us-east-1"
+ revoke_rules_on_delete = false
+ tags              = {
+   + "Name" = "dev-web-sg-0"
}
+ tags_all          = {
+   + "Name" = "dev-web-sg-0"
}
+ vpc_id            = "vpc-0a9005120b91ce98c"
}

Plan: 3 to add, 0 to change, 0 to destroy.

Changes to Outputs:
+ instance_id      = (known after apply)
+ webserver_public_ip = (known after apply)
module.myapp-webserver.aws_key_pair.ssh-key: Creating...
module.myapp-webserver.aws_security_group.web_sg: Creating...
module.myapp-webserver.aws_key_pair.ssh-key: Creation complete after 1s [id=dev-serverkey-0]
module.myapp-webserver.aws_security_group.web_sg: Creation complete after 5s [id=sg-03cf0ef7bb487c907]
module.myapp-webserver.aws_instance.myapp-server: Creating...
^[[6~[[A^[[Amodule.myapp-webserver.aws_instance.myapp-server: Still creating... [00m10s elapsed]
module.myapp-webserver.aws_instance.myapp-server: Creation complete after 15s [id=i-0de4cf3723d982a6f]

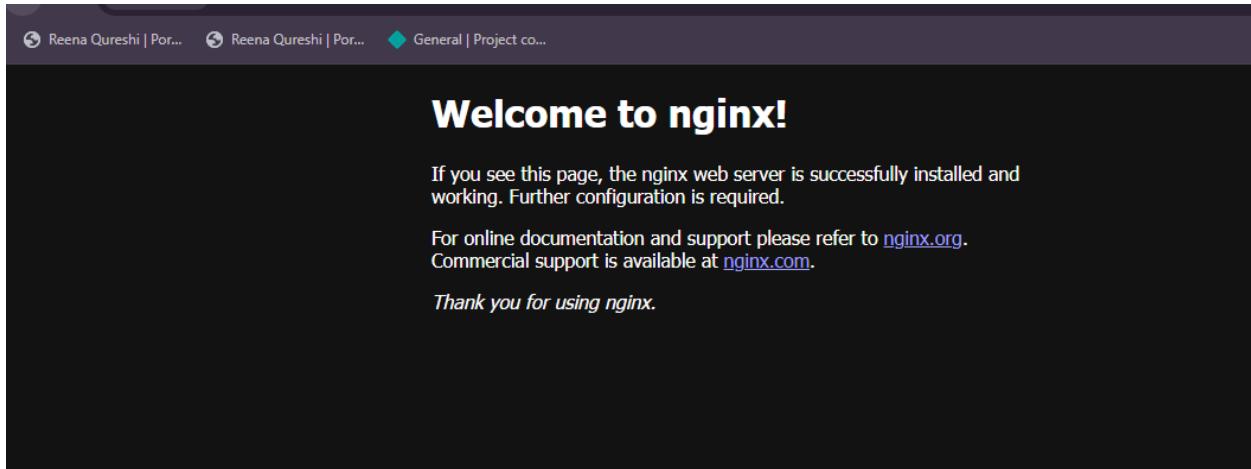
Apply complete! Resources: 3 added, 0 changed, 0 destroyed.

Outputs:
instance_id = "i-0de4cf3723d982a6f"
```

## OutPut:

```
greenaquareshi ~ /Lab12 $ cat screenshots/tasks_5_n
Task 5 Instance IP: 13.219.245.239
Testing Nginx with webserver module...
```

Browser test:



Destroying Resources:

```
Enter a value: yes

module.myapp-webserver.aws_instance.myapp-server: Destroying... [id=i-0de4cf3723d982a6f]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-0de4cf3723d982a6f, 00m1s]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-0de4cf3723d982a6f, 00m2s]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-0de4cf3723d982a6f, 00m3s]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-0de4cf3723d982a6f, 00m4s]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-0de4cf3723d982a6f, 00m5s]
module.myapp-webserver.aws_instance.myapp-server: Destruction complete after 52s
module.myapp-webserver.aws_key_pair.ssh-key: Destroying... [id=dev-serverkey-0]
module.myapp-webserver.aws_security_group.web_sg: Destroying... [id=sg-03cf0ef7bb487c907]
module.myapp-webserver.aws_key_pair.ssh-key: Destruction complete after 0s
module.myapp-webserver.aws_security_group.web_sg: Destruction complete after 2s
```

## Task 6 — Configure HTTPS with self-signed certificates

In this task, you will configure Nginx to serve traffic over HTTPS using self-signed certificates.

Update `entry-script.sh` with SSL configuration:

```
# Update and install required @reenaqureshi ② ~/Lab12 $ cat > entry-script.sh << 'EOF'
> #!/bin/bash
> set -e
>
> # Update and install required packages
> sudo yum update -y
> sudo amazon-linux-extras enable nginx1
> sudo yum install -y nginx openssl
>
> # Create directories for SSL certificates
> sudo mkdir -p /etc/ssl/private
> sudo mkdir -p /etc/ssl/certs
>
> # Get IMDSv2 token
> TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" \
>   -H "X-aws-ec2-metadata-token-ttl-seconds: 21600")
>
> # Get current public IP and hostname
> PUBLIC_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
```

Apply the configuration:

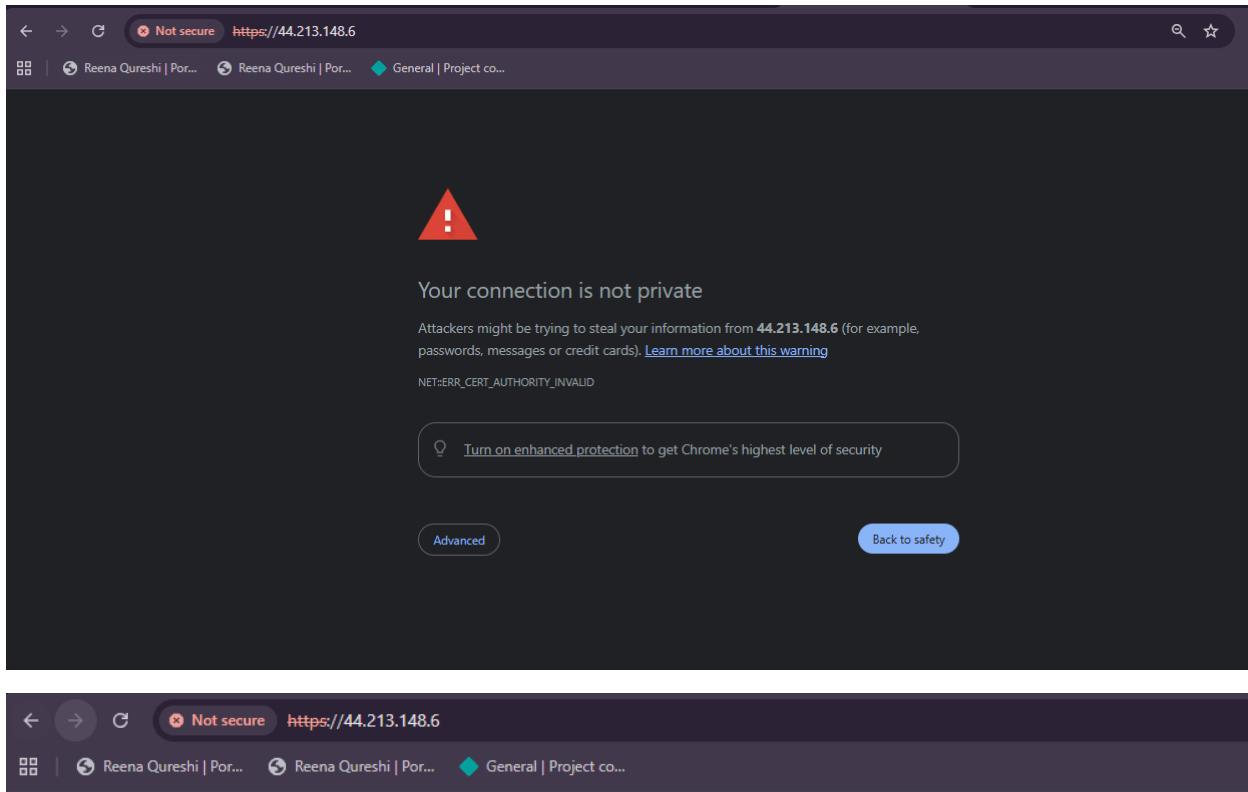
```
terraform apply -auto-approve
Accept security warning for self-signed certificate
@reenaqureshi ② ~/Lab12 $ terraform apply
data.http.my_ip: Reading...
data.http.my_ip: Read complete after 0s [id=https://icanhazip.com]
module.myapp-subnet.data.aws_vpc.default: Reading...
module.myapp-webserver.aws_key_pair.ssh-key: Refreshing state... [id=dev-serverkey-0]
module.myapp-subnet.data.aws_vpc.default: Read complete after 2s [id=vpc-0a9005120b91ce98c]
module.myapp-subnet.data.aws_subnet.default: Reading...
module.myapp-webserver.aws_security_group.web_sg: Refreshing state... [id=sg-0d646d1910f202a86]
module.myapp-subnet.data.aws_subnet.default: Read complete after 1s [id=subnet-00447af4668f7212a]
module.myapp-webserver.aws_instance.myapp-server: Refreshing state... [id=i-04e12a68150b3fa35]

No changes. Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes will be made.
```

Outputs:

```
instance_id = "i-04f7da1b82a4ba3cd"
webserver_public_ip = "44.213.148.6"
```



## “ HTTPS Server is Working!

### Server Information:

**Status:** HTTPS Active

**HTTP:** Redirects to HTTPS

**SSL:** Self-signed certificate installed

**Time:** \$(date)

**IP:** \$(curl -s http://169.254.169.254/latest/meta-data/public-ipv4)

### Test Results:

“ HTTP â†’ HTTPS redirect working

“ SSL certificate configured

“ Nginx serving content

## Task 7 — Configure Nginx as reverse proxy

Create apache.sh script for backend web server:

Apply the configuration:

```
terraform apply -auto-approve
```

```
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-04f7da1b82a4ba3cd, 00m50s elapsed]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-04f7da1b82a4ba3cd, 01m00s elapsed]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-04f7da1b82a4ba3cd, 01m10s elapsed]
module.myapp-webserver.aws_instance.myapp-server: Still destroying... [id=i-04f7da1b82a4ba3cd, 01m20s elapsed]
module.myapp-webserver.aws_instance.myapp-server: Destruction complete after 1m23s
module.myapp-webserver.aws_key_pair.ssh-key: Destroying... [id=dev-https-key-https]
module.myapp-webserver.aws_security_group.web_sg: Destroying... [id=sg-07c3aff118e0b1f17]
module.myapp-webserver.aws_key_pair.ssh-key: Destruction complete after 1s
module.myapp-webserver.aws_security_group.web_sg: Destruction complete after 2s

Apply complete! Resources: 9 added, 0 changed, 3 destroyed.

Outputs:

all_ips = {
  "nginx_proxy" = "44.213.128.103"
  "web1" = "44.222.174.52"
  "web2" = "44.213.125.100"
}
nginx_proxy_instance_id = "i-00be26f0e8b4d3268"
nginx_proxy_public_ip = "44.213.128.103"
web1_instance_id = "i-00b59f4be09b7614f"
web1_public_ip = "44.222.174.52"
web2_instance_id = "i-00b904390a0e585d3"
web2_public_ip = "44.213.125.100"
@reenaquireshi ~ ~/Lab12 $ 
@reenaquireshi ~ ~/Lab12 $ # Get all IPs
@reenaquireshi ~ ~/Lab12 $ terraform output
all_ips = {
  "nginx_proxy" = "44.213.128.103"
  "web1" = "44.222.174.52"
}
nginx_proxy_instance_id = "i-00be26f0e8b4d3268"
nginx_proxy_public_ip = "44.213.128.103"
web1_instance_id = "i-00b59f4be09b7614f"
web1_public_ip = "44.222.174.52"
web2_instance_id = "i-00b904390a0e585d3"
web2_public_ip = "44.213.125.100"
@reenaquireshi ~ ~/Lab12 $
```

Get the outputs:

```
terraform output
```

```
@reenaquireshi ~ ~/Lab12 $ # Get all IPs
@reenaquireshi ~ ~/Lab12 $ terraform output
all_ips = {
  "nginx_proxy" = "44.213.128.103"
  "web1" = "44.222.174.52"
  "web2" = "44.213.125.100"
}
nginx_proxy_instance_id = "i-00be26f0e8b4d3268"
nginx_proxy_public_ip = "44.213.128.103"
web1_instance_id = "i-00b59f4be09b7614f"
web1_public_ip = "44.222.174.52"
web2_instance_id = "i-00b904390a0e585d3"
web2_public_ip = "44.213.125.100"
@reenaquireshi ~ ~/Lab12 $
```

**SSH into the webserver (Nginx proxy server)**

```
@greenaquareishi ~ ~/Lab12 $ ssh ec2-user@44.222.174.52
The authenticity of host '44.222.174.52 (44.222.174.52)' can't be established.
ED25519 key fingerprint is SHA256:ySsSyHwpyDRYYohcNQ13XoSjMwqH7T817sBwWbCxA2Y.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '44.222.174.52' (ED25519) to the list of known hosts.

          #
~\_ #####      Amazon Linux 2
~~ \#####\
~~ \###|      AL2 End of Life is 2026-06-30.
~~  \#/      V~'__->
~~   /      A newer version of Amazon Linux is available!
~~_.-/      Amazon Linux 2023, GA and supported until 2028-03-15.
~/m/`-`/      https://aws.amazon.com/linux/amazon-linux-2023/
[ec2-user@backend-server-00b5 ~]$
```

Edit the Nginx configuration:

```
sudo vim /etc/nginx/nginx.conf
```

```
# TASK 7. Reverse proxy to web1
location / {
    proxy_pass http://44.213.125.100:80;  # WEB1_IP
    proxy_set_header Host \$host;
    proxy_set_header X-Real-IP \$remote_addr;
    proxy_set_header X-Forwarded-For \$proxy_add_x_forwarded_for;
    proxy_set_header X-Forwarded-Proto \$scheme;
}
}

NGINX
```

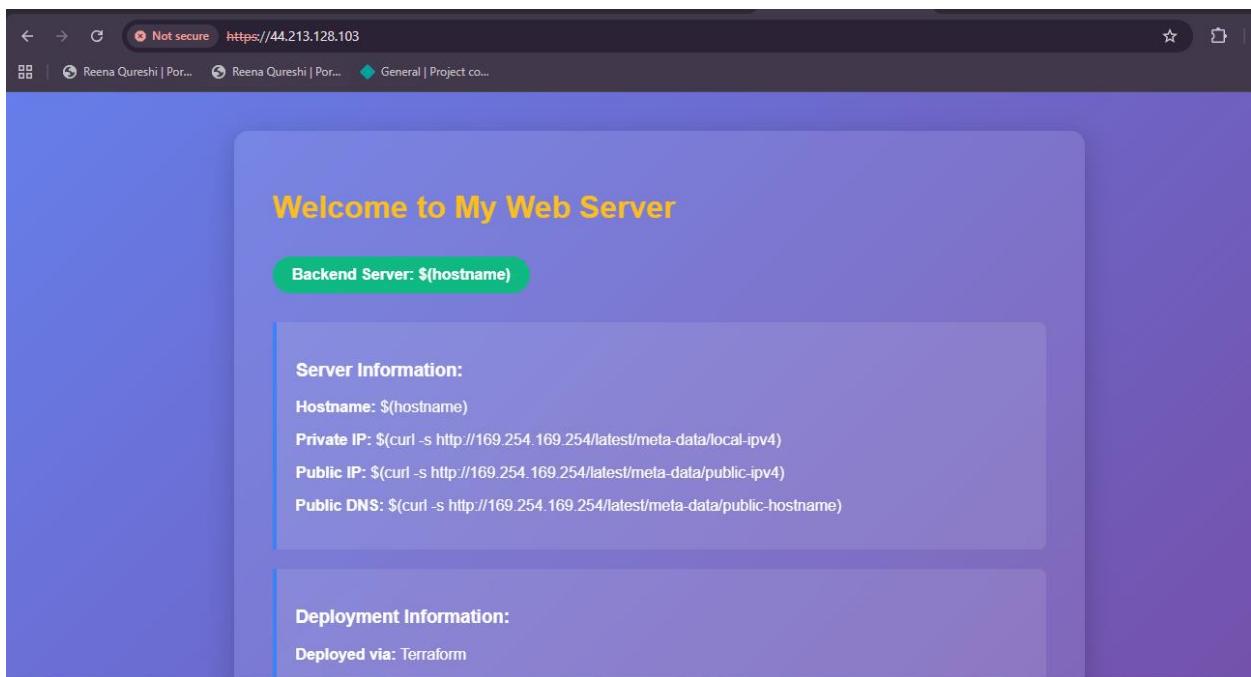
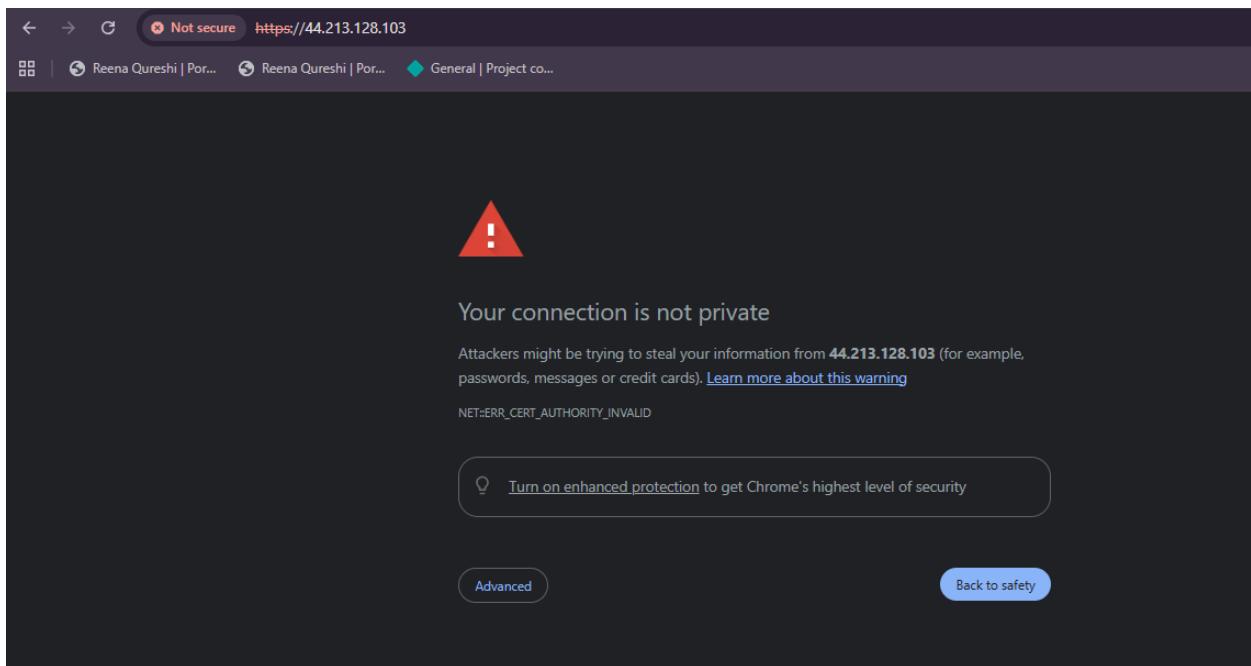
Restart Nginx:

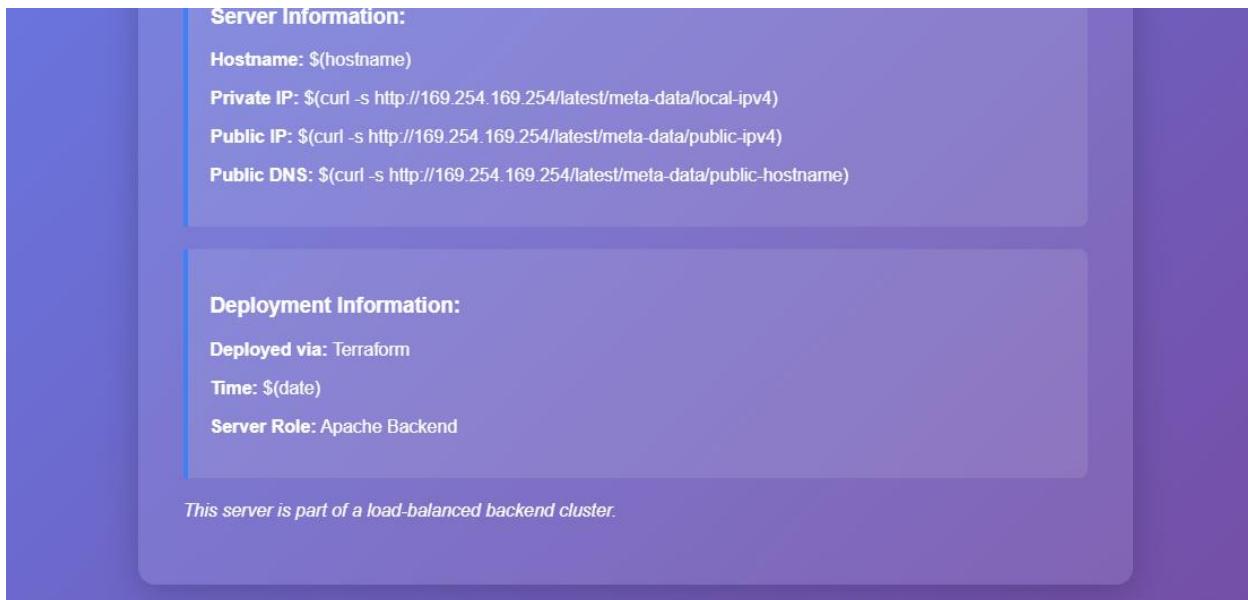
```
sudo systemctl restart nginx
```

```
[ec2-user@ip-172-31-8-106 ~]$ # Restart nginx
[ec2-user@ip-172-31-8-106 ~]$ sudo systemctl restart nginx
[ec2-user@ip-172-31-8-106 ~]$
```

View Nginx logs and configuration files:

```
cat /var/log/nginx/error.log
```





## Task 8 — Configure Nginx as load balancer

In this task, you will add a second backend server and configure Nginx

Add the second web server module to `main.tf`:

Update `outputs.tf`:

Apply the configuration:

```
module.nginx-proxy.aws_security_group.web_sg: Creation complete after 5s [id=sg-0c9d0e2df7a355f4c]
module.myapp-web-2.aws_instance.myapp-server: Creating...
module.nginx-proxy.aws_instance.myapp-server: Creating...
module.myapp-web-1.aws_instance.myapp-server: Still creating... [00m10s elapsed]
module.myapp-web-2.aws_instance.myapp-server: Still creating... [00m10s elapsed]
module.nginx-proxy.aws_instance.myapp-server: Still creating... [00m10s elapsed]
module.myapp-web-1.aws_instance.myapp-server: Still creating... [00m20s elapsed]
module.myapp-web-2.aws_instance.myapp-server: Still creating... [00m20s elapsed]
module.nginx-proxy.aws_instance.myapp-server: Still creating... [00m20s elapsed]
module.nginx-proxy.aws_instance.myapp-server: Creation complete after 24s [id=i-0656851113f46f24c]
module.myapp-web-1.aws_instance.myapp-server: Still creating... [00m30s elapsed]
module.myapp-web-2.aws_instance.myapp-server: Still creating... [00m30s elapsed]
module.myapp-web-2.aws_instance.myapp-server: Creation complete after 34s [id=i-04b86cb564f01ae0a]
module.myapp-web-1.aws_instance.myapp-server: Creation complete after 34s [id=i-0a32a30c6aad1ee29]
```

Get all outputs:

```
all_servers = {
  "backend_1" = "3.233.237.119"
  "backend_2" = "13.220.9.105"
  "nginx_proxy" = "34.239.93.60"
}
nginx_proxy_public_ip = "34.239.93.60"
web1_public_ip = "3.233.237.119"
web2_public_ip = "13.220.9.105"
@greenacuposhi ~ % /Lab12 $
```



Reena Qureshi | Por...



Reena Qureshi | Por...



General | Project co...

## Nginx Proxy Server

**Status:** Ready for Tasks 7-10 Configuration

**IP:** 3.233.237.119

**Tasks:**

- Task 7: Reverse Proxy
- Task 8: Load Balancer
- Task 9: High Availability
- Task 10: Caching

Configure via: /etc/nginx/nginx.conf



Reena Qureshi | Por...



Reena Qureshi | Por...



General | Project co...

## Nginx Proxy Server

**Status:** Ready for Tasks 7-10 Configuration

**IP:** 13.220.9.105

**Tasks:**

- Task 7: Reverse Proxy
- Task 8: Load Balancer
- Task 9: High Availability
- Task 10: Caching

Configure via: /etc/nginx/nginx.conf

## TASK 9: Configure high availability with backup servers

```
@reenaqureshi ~ ~/Lab12 $ ssh ec2-user@34.239.93.60
The authenticity of host '34.239.93.60 (34.239.93.60)' can't be established.
ED25519 key fingerprint is SHA256:9/qHZCWWy2xTWWiHFZR1G2GFTfnhO30qbbTwNeSavjs.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '34.239.93.60' (ED25519) to the list of known hosts.

      _#
  ~\_\_ #####_      Amazon Linux 2
  ~~ \_\#\#\#\#\_ 
  ~~ \#\#\#|      AL2 End of Life is 2026-06-30.
  ~~ \#/ \_-
  ~~ V~' '-'>
  ~~ /      A newer version of Amazon Linux is available!
  ~~ . /      Amazon Linux 2023, GA and supported until 2028-03-15.
  _/m/      https://aws.amazon.com/linux/amazon-linux-2023/
[ec2-user@ip-172-31-14-17 ~]$
```

```
[ec2-user@ip-172-31-14-17 ~]$ systemctl status nginx
● nginx.service - The nginx HTTP and reverse proxy server
   Loaded: loaded (/usr/lib/systemd/system/nginx.service; enabled; vendor preset: disabled)
     Active: active (running) since Sun 2026-01-04 13:48:21 UTC; 32min ago
       Process: 6419 ExecStart=/usr/sbin/nginx (code=exited, status=0/SUCCESS)
      Process: 6417 ExecStartPre=/usr/sbin/nginx -t (code=exited, status=0/SUCCESS)
      Process: 6415 ExecStartPre=/usr/bin/rm -f /run/nginx.pid (code=exited, status=0/SUCCESS)
    Main PID: 6422 (nginx)
      CGroup: /system.slice/nginx.service
              ├─6422 nginx: master process /usr/sbin/nginx
              ├─6425 nginx: worker process
              └─6426 nginx: cache manager process

Jan 04 13:48:22 ip-172-31-14-17.ec2.internal nginx[6417]: nginx: the configuration file /etc/nginx/nginx.conf
Jan 04 13:48:22 ip-172-31-14-17.ec2.internal nginx[6417]: nginx: configuration file /etc/nginx/nginx.conf
Jan 04 13:48:21 ip-172-31-14-17.ec2.internal systemd[1]: Starting The nginx HTTP and reverse proxy server
Jan 04 13:48:21 ip-172-31-14-17.ec2.internal systemd[1]: Started The nginx HTTP and reverse proxy server.
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-31-14-17 ~]$
```

Server running ON web 2:



Reena Qureshi | Por...



Reena Qureshi | Por...



General | Project co...

## Nginx Proxy Server

**Status:** Ready for Tasks 7-10 Configuration

**IP:** 13.220.9.105

### Tasks:

- Task 7: Reverse Proxy
- Task 8: Load Balancer
- Task 9: High Availability
- Task 10: Caching

Configure via: /etc/nginx/nginx.conf

## Task 10 — Enable Nginx caching

### 1. SSH into the webserver:

```
ssh ec2-user@<webserver-public-ip>
```

```
@reenaquireshi ~ ~/Lab12 $ ssh ec2-user@34.239.93.60
The authenticity of host '34.239.93.60 (34.239.93.60)' can't be established.
ED25519 key fingerprint is SHA256:9/qHZCWWy2xTWWiHFZR1G2GFTfnhO30qbbTwNeSavjs.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '34.239.93.60' (ED25519) to the list of known hosts.

      #_
      ~\_ #####_          Amazon Linux 2
      ~~ \#####\_
      ~~  \###|          AL2 End of Life is 2026-06-30.
      ~~   \|/_____
      ~~    V~' '-'>
      ~~~   /_____
      ~~~  /_/_/
      _/m/_/ A newer version of Amazon Linux is available!
                                         Amazon Linux 2023, GA and supported until 2028-03-15.
                                         https://aws.amazon.com/linux/amazon-linux-2023/
[ec2-user@in-172-31-14-17 ~]$
```